

malising flows are a family of distributions defined by a base distribution (usually a standard Gaussian), and a *diffeomorphism*  $f$ , *i.e.* a differentiable bijection with a differentiable inverse. Specifically, we set  $p_\theta(x(t+\tau)|x(t))$  as the density of the distribution defined by the following generative process:

$$\begin{aligned} z^p, z^v &\sim \mathcal{N}(0, I) \\ x^p(t+\tau), x^v(t+\tau) &:= f_\theta(z^p, z^v; x^p(t), x^v(t)). \end{aligned} \quad (3)$$

Here  $z^p \in \mathbb{R}^{3N}$  and  $z^v \in \mathbb{R}^{3N}$ . For all settings of  $\theta$  and  $x(t)$ ,  $f_\theta(\cdot; x(t))$  is a diffeomorphism that takes the latent variables  $(z^p, z^v) \in \mathbb{R}^{6N}$  to  $(x^p(t+\tau), x^v(t+\tau)) \in \mathbb{R}^{6N}$ . The conditioning state  $x(t)$  parameterises a family of diffeomorphisms, defining a *conditional* normalising flow (Winkler et al., 2019). Note that there are no invertibility constraints on the mapping from the conditioning state  $x(t)$  to the output  $x(t+\tau)$ , only the map from  $z$  to  $x(t+\tau)$  must be invertible. Using the change of variables formula, we can evaluate  $p_\theta(x(t+\tau)|x(t))$  analytically as:

$$\mathcal{N}(f_\theta^{-1}(x(t+\tau); x(t)); 0, I) \left| \det \mathcal{J}_{f_\theta^{-1}(\cdot; x(t))}(x(t+\tau)) \right|,$$

where  $f_\theta^{-1}(\cdot; x(t)) : \mathbb{R}^{6N} \rightarrow \mathbb{R}^{6N}$  is the inverse of the diffeomorphism  $f_\theta(\cdot; x(t))$ , and  $\mathcal{J}_{f_\theta^{-1}(\cdot; x(t))}(x(t+\tau))$  denotes the Jacobian of  $f_\theta^{-1}(\cdot; x(t))$  evaluated at  $x(t+\tau)$ .

### 3.2. Dataset generation

We now describe the dataset used to train the flow. We generate MD trajectories by integrating Equation (2) using the *OpenMM* library (Eastman et al., 2017). We simulate small proteins (peptides) in implicit water, *i.e.*, without explicitly modelling the degrees of freedom of the water molecules. Specifically, we generate a dataset of trajectories  $\mathcal{D} = \{\mathcal{T}_i\}_{i=1}^P$ , where  $P$  is the number of peptides. For each peptide  $i$ , we generate an MD trajectory that is temporally sampled with a spacing of  $\tau$ , so that  $\mathcal{T}_i = (x(0), x(\tau), x(2\tau), \dots)$ . During training, we randomly sample pairs  $x(t), x(t+\tau)$  from  $\mathcal{D}$ . Each pair represents a sample from the conditional distribution  $\mu(x(t+\tau)|x(t))$ . These samples are used as examples to train the parameters  $\theta$  of the flow. Additional details are provided in Appendix D. Since the flow is trained on trajectory data from *multiple* peptides, we can deploy it at test time to generalise to *new* peptides not seen in the training data.

### 3.3. Augmented normalising flows

Typically in molecular simulations, we are primarily interested in the distribution of the *positions*  $x^p$ , rather than the velocities  $x^v$ . Thus, it is not necessary for  $x^v(t), x^v(t+\tau)$  to represent the actual velocities of the atoms in Equation (3). We hence simplify the learning problem by treating  $x^v$  as

*non-physical auxiliary variables* within the *augmented normalising flow* framework (Huang et al., 2020). For each datapoint  $x(t) = x^p(t), x^v(t)$  in  $\mathcal{D}$ , instead of obtaining  $x^v(t)$  by recording the velocities in the MD trajectory, we *discard* the MD velocity and independently draw  $x^v(t) \sim \mathcal{N}(0, I)$ . The auxiliary variables  $x^v(t)$  now contain no information about the future state  $x^p(t+\tau), x^v(t+\tau)$ , since  $x^v(t)$  and  $x^v(t+\tau)$  are drawn independently. Hence we can simplify  $f_\theta$  to depend only on  $x^p(t)$ , with  $x^p(t+\tau), x^v(t+\tau) := f_\theta(z^p, z^v; x^p(t))$ . This raises the question of why auxiliary variables are necessary: we could instead directly model  $p_\theta(x^p(t+\tau)|x^p(t))$ , without the need for  $x^v$ . We include auxiliary variables for two reasons: First, they increase the expressivity of the distribution for  $x^p$  without a prohibitive increase in computational cost (Huang et al., 2020; Chen et al., 2020). Second, constructing a conditional flow that respects *permutation equivariance* is simplified with auxiliary variables — we discuss this in more detail in Section 4.1.

### 3.4. Targeting the Boltzmann distribution with MCMC

Once the flow  $p_\theta(x(t+\tau)|x(t))$  has been trained, we use it as a proposal distribution in an MCMC method to target the joint distribution of the positions  $x^p$  and the auxiliary variables  $x^v$ , which has density:

$$\mu_{\text{aug}}(x^p, x^v) \propto \exp\left(-\frac{U(x^p)}{k_B T}\right) \mathcal{N}(x^v; 0, I). \quad (4)$$

Let  $m$  index the states in the Markov chain. Starting from an initial state  $X_0 = (X_0^p, X_0^v) \in \mathbb{R}^{6N}$  at  $m = 0$ , we iterate:

$$\tilde{X}_m \sim p_\theta(\cdot | X_m^p) \quad (5)$$

$$X_{m+1} := \begin{cases} \tilde{X}_m & \text{with probability } \alpha(X_m, \tilde{X}_m) \\ X_m & \text{with probability } 1 - \alpha(X_m, \tilde{X}_m) \end{cases} \quad (6)$$

where  $\alpha(X_m, \tilde{X}_m)$  is the *Metropolis-Hastings (MH) acceptance ratio* (Metropolis et al., 1953) targeting Equation (4):

$$\alpha(X, \tilde{X}) = \min\left(1, \frac{\mu_{\text{aug}}(\tilde{X})p_\theta(X | \tilde{X}^p)}{\mu_{\text{aug}}(X)p_\theta(\tilde{X} | X^p)}\right) \quad (7)$$

The flow used for  $p_\theta$  must allow for efficient sampling *and* exact likelihood evaluation, which is crucial for fast implementation of Equations (5) and (7). Additionally, after each MH step, we resample the auxiliary variables  $X^v$  using a *Gibbs sampling* update:

$$(X_m^p, X_m^v) \leftarrow (X_m^p, \epsilon), \quad \epsilon \sim \mathcal{N}(0, I). \quad (8)$$

Iterating these updates yields a sample  $X_m^p, X_m^v \sim \mu_{\text{aug}}$  as  $m \rightarrow \infty$ . To obtain a Boltzmann-distributed sample of the positions  $X_m^p \sim \mu$ , we simply discard the auxiliary variables  $X_m^v$ . In practice, sending  $m \rightarrow \infty$  is infeasible. Instead, we fix a computational budget and simulate the chain until the budget is reached. Algorithm 1 shows pseudocode for the MCMC procedure, using a *batch sampling* procedure for the proposals which significantly speeds up sampling.